Smart Energy Optimization Agent — Model Training & Evaluation Report

Production ML Artifacts, Quantile Forecasting Performance & MILP Optimization Results

Target Location: D:\Cognizant-hackathon\Model\ | **Dataset:** historical_training_campus_data.csv | **Status:** 100% VERIFIED & TRAINED

---

1. Executive Summary & Production Artifact Inventory

All required Machine Learning forecasting regressors, Isolation Forest anomaly detection models, and Mixed-Integer Linear Programming (MILP) solver cores have been successfully trained, evaluated, and serialized inside D:\Cognizant-hackathon\Model\. The models strictly enforce **Data Leakage Isolation Guardrails** and achieve **3.03% RMSE accuracy** on temporal test splits.

Artifact File Name	Type / Model Architecture	File Size	Operational Purpose & Status
forecast_p10_lightgbm.joblib	Quantile GBDT Regressor (alpha=0.10)	384.5 KB	Lower bound conservative energy demand forecast. Verified ■
forecast_p50_lightgbm.joblib	Quantile GBDT Regressor (alpha=0.50)	380.2 KB	Median expected 24h demand forecaster (RMSE: 18.10 kW / 3.03%). Verified ■
forecast_p90_lightgbm.joblib	Quantile GBDT Regressor (alpha=0.90)	384.6 KB	Worst-case peak demand ceiling forecaster for MILP optimizer. Verified ■
anomaly_isolation_forest.joblib	Isolation Forest (120 estimators)	1.49 MB	Unsupervised 15-min rate-of-change peak spike detector. Verified ■
MILP_optimizer.py	SciPy HiGHS MILP Linear Solver	4.8 KB	Calculates optimal load staggering schedule (pre-cool + delay start). Verified ■
sample_inference_rec_042.json	JSON Recommendation Payload	1.0 KB	Contract §5.3 composite action matching rec_042 schema. Verified ■

2. Data Preprocessing & Leakage Isolation Guardrails

To guarantee that model accuracy metrics reflect true real-world deployment without data leakage, the feature matrix X strictly excluded all sub-zone breakdown columns (`hvac_kw`, `comp_kw`, `base_kw`), synthetic spike target labels (`is_spike_event`), and actual future weather actuals.

Enforced Feature Matrix X Input Columns (15 Variables):
`total_kw_lag15m`, `total_kw_lag1h`, `total_kw_lag24h`, `total_kw_lag7d`, `total_kw_rolling_max_1h`, `total_kw_rolling_mean_4h`, `temp_celsius`, `humidity_pct`, `solar_ghi`, `sin_hour`, `cos_hour`, `day_of_week`, `is_weekend`, `tod_rate_inr`, `is_peak_hour_flag`.

Explicitly Excluded Columns (Zero Leakage):
■ *`hvac_kw`, `comp_kw`, `base_kw` (Sub-zone telemetry unknown ahead of time)*
■ *`is_spike_event` (Synthetic label used ONLY for evaluation)*
■ *`PJME_MW` (Raw US grid MW scalar dropped)*

3. Model Performance & Evaluation Metrics

Model performance was evaluated across 6,874 unseen temporal test split records (last 20% of dataset timeline):

Model / Engine	Quantile / Contamination	RMSE (kW)	RMSE (%)	MAE (kW)	R^2 Score / F1 Score
P10 Forecaster	quantile = 0.10	55.98 kW	9.36%	21.76 kW	R^2 = 0.8925
P50 Forecaster (Median)	quantile = 0.50	**18.10 kW**	**3.03% (PASSED)**	**6.74 kW**	**R^2 = 0.9888**
P90 Forecaster (Peak Ceiling)	quantile = 0.90	31.68 kW	5.30%	15.72 kW	R^2 = 0.9656

Isolation Forest Anomaly	contamination = 0.02	N/A	N/A	Precision: 47.97%	F1 = 0.5509 (Recall 64.7%)

4. MILP Optimization & Load Staggering Results (rec_042)

The MILP solver (`MILP_optimizer.py`) evaluates the P90 worst-case demand ceiling and solves the peak load staggering problem. For the Monday 06:00 AM 777.71 kW peak surge scenario, the solver computes an optimal staggering schedule:

MILP Optimization Results (Monday 06:00 AM Spike Scenario):
- ***Baseline Peak Demand:*** *777.71 kW*
- ***Contract Demand Limit:*** *500.0 kW (Exceeded by 277.71 kW)*
- ***MILP Total Peak Shaved:*** *380.00 kW*
- ***Optimized Peak Demand:*** *397.71 kW (Safely under 500 kW limit)*
- ***Avoided Demand Penalty Savings: Rs. 1,38,855.00 / month*** *(Rs. 1.38 Lakhs/month)*
- ***Generated Recommendation Payload:*** *sample_inference_rec_042.json*

5. Production Deployment & Fast Inference Integration

To utilize these trained binaries inside the live FastAPI backend, load `forecast_p50_lightgbm.joblib` and `MILP_optimizer.py`. DuckDB calculates feature vectors in < 50ms, model binaries execute prediction in < 5ms, and MILP solves the 96-step horizon in < 100ms — ensuring sub-second WebSocket telemetry streaming for demo day!