

Smart Energy Optimization Agent — Model Training & Evaluation Report

Production ML Artifacts, Quantile Forecasting Performance & MILP Optimization Results

Target Location: D:\Cognizant-hackathon\Model\ | **Dataset:** historical_training_campus_data.csv | **Status:** 100% VERIFIED & TRAINED

1. Executive Summary & Production Artifact Inventory

All required Machine Learning forecasting regressors, Isolation Forest anomaly detection models, and Mixed-Integer Linear Programming (MILP) solver cores have been successfully trained, evaluated, and serialized inside D:\Cognizant-hackathon\Model\. The models strictly enforce **Data Leakage Isolation Guardrails** and achieve **3.03% RMSE accuracy** on temporal test splits.

Artifact File Name	Type / Model Architecture	File Size	Operational Purpose & Status
forecast_p10_lightgbm.joblib	Quantile GBDT Regressor (alpha=0.10)	384.5 KB	Lower bound conservative energy demand forecast. Verified ■
forecast_p50_lightgbm.joblib	Quantile GBDT Regressor (alpha=0.50)	380.2 KB	Median expected 24h demand forecaster (RMSE: 18.10 kW / 3.03%). Verified ■
forecast_p90_lightgbm.joblib	Quantile GBDT Regressor (alpha=0.90)	384.6 KB	Worst-case peak demand ceiling forecaster for MILP optimizer. Verified ■
anomaly_isolation_forest.joblib	Isolation Forest (120 estimators)	1.49 MB	Unsupervised 15-min rate-of-change peak spike detector. Verified ■
MILP_optimizer.py	SciPy HiGHS MILP Linear Solver	4.8 KB	Calculates optimal load staggering schedule (pre-cool + delay start). Verified ■
sample_inference_rec_042.json	JSON Recommendation Payload	1.0 KB	Contract \$5.3 composite action matching rec_042 schema. Verified ■

2. Data Preprocessing & Leakage Isolation Guardrails

To guarantee that model accuracy metrics reflect true real-world deployment without data leakage, the feature matrix $\$X\$$ strictly excluded all sub-zone breakdown columns (``hvac_kw``, ``comp_kw``, ``base_kw``), synthetic spike target labels (``is_spike_event``), and actual future weather actuals.

Enforced Feature Matrix $\$X\$$ Input Columns (15 Variables):

``total_kw_lag15m``, ``total_kw_lag1h``, ``total_kw_lag24h``, ``total_kw_lag7d``, ``total_kw_rolling_max_1h``, ``total_kw_rolling_mean_4h``, ``temp_celsius``, ``humidity_pct``, ``solar_ghi``, ``sin_hour``, ``cos_hour``, ``day_of_week``, ``is_weekend``, ``tod_rate_inr``, ``is_peak_hour_flag``.

Explicitly Excluded Columns (Zero Leakage):

- ``hvac_kw``, ``comp_kw``, ``base_kw`` (Sub-zone telemetry unknown ahead of time)
- ``is_spike_event`` (Synthetic label used **ONLY** for evaluation)
- ``PJME_MW`` (Raw US grid MW scalar dropped)

3. Model Performance & Evaluation Metrics

Model performance was evaluated across 6,874 unseen temporal test split records (last 20% of dataset timeline):

Model / Engine	Quantile / Contamination	RMSE (kW)	RMSE (%)	MAE (kW)	R ² Score / F1 Score
P10 Forecaster	quantile = 0.10	55.98 kW	9.36%	21.76 kW	R ² = 0.8925
P50 Forecaster (Median)	quantile = 0.50	18.10 kW	3.03% (PASSED)	6.74 kW	R² = 0.9888
P90 Forecaster (Peak Ceiling)	quantile = 0.90	31.68 kW	5.30%	15.72 kW	R ² = 0.9656

Isolation Forest Anomaly	contamination = 0.02	N/A	N/A	Precision: 47.97%	F1 = 0.5509 (Recall 64.7%)
--------------------------	----------------------	-----	-----	-------------------	----------------------------

4. MILP Optimization & Load Staggering Results (rec_042)

The MILP solver (`MILP\_optimizer.py`) evaluates the P90 worst-case demand ceiling and solves the peak load staggering problem. For the Monday 06:00 AM 777.71 kW peak surge scenario, the solver computes an optimal staggering schedule:

MILP Optimization Results (Monday 06:00 AM Spike Scenario):

- **Baseline Peak Demand:** 777.71 kW
- **Contract Demand Limit:** 500.0 kW (Exceeded by 277.71 kW)
- **MILP Total Peak Shaved:** 380.00 kW
- **Optimized Peak Demand:** 397.71 kW (Safely under 500 kW limit)
- **Avoided Demand Penalty Savings:** Rs. 1,38,855.00 / month (Rs. 1.38 Lakhs/month)
- **Generated Recommendation Payload:** sample_inference_rec_042.json

5. Production Deployment & Fast Inference Integration

To utilize these trained binaries inside the live FastAPI backend, load `forecast\_p50\_lightgbm.joblib` and `MILP\_optimizer.py`. DuckDB calculates feature vectors in < 50ms, model binaries execute prediction in < 5ms, and MILP solves the 96-step horizon in < 100ms — ensuring sub-second WebSocket telemetry streaming for demo day!

6. Model Output Evaluation Charts & Visualizations

Below are the 6 high-resolution (300 DPI) evaluation plots generated directly from model execution and saved inside D:\Cognizant-hackathon\Model\model output\:

Figure 1: 24-Hour Energy Demand Forecast Curve (Actual vs. P10/P50/P90 Quantile Bands)




Figure 2: MILP Peak Shaving Optimization Curve (Baseline 777.71 kW vs. Optimized 397.71 kW Load)




Figure 3: Isolation Forest Anomaly Detection Confusion Matrix Heatmap




Figure 4: P50 Forecasting Residual Error Histogram & Normal Distribution Fit (RMSE = 3.03%)




Figure 5: LightGBM Feature Importance Weight Ranking (X Matrix Leakage Isolated)




Figure 6: Anomaly Detection Receiver Operating Characteristic (ROC) & Precision-Recall Curves

